

# HW1

28 февраля 2017 г.

## 1 Домашнее задание 1

1.1 Гаркавый Андрей, 494 группа

1.1.1 1. Метод k ближайших соседей

```
In [4]: from sklearn import cross_validation, datasets, metrics, neighbors, model_selection
        from matplotlib.colors import ListedColormap
        import numpy as np
        %pylab inline
```

```
/usr/local/lib/python3.5/dist-packages/sklearn/cross_validation.py:44: DeprecationWarning: This
    "This module will be removed in 0.20.", DeprecationWarning)
```


Populating the interactive namespace from numpy and matplotlib

```
In [5]: colors = ListedColormap(['red', 'blue', 'purple', 'green'])
        light_colors = ListedColormap(['lightcoral', 'lightblue', 'pink', 'lightgreen'])

        def plot_2d_dataset(data, colors):
            pyplot.figure(figsize(8, 8))
            pyplot.scatter(list(map(lambda x: x[0], data[0])), list(map(lambda x: x[1], data[0]))

In [6]: classification_problem = datasets.make_classification(n_samples=2000, n_features=2, n_in
                                                    n_classes=4, n_redundant=0, n_clus
                                                    random_state=1)

        plot_2d_dataset(classification_problem, colors)
```



```
In [7]: train_data, test_data, train_labels, test_labels = cross_validation.train_test_split(cla
cla
tes
ran

In [8]: estimator = neighbors.KNeighborsClassifier()
estimator.fit(train_data, train_labels)

Out[8]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
metric_params=None, n_jobs=1, n_neighbors=5, p=2,
weights='uniform')

In [9]: predictions = estimator.predict(test_data)
metrics.accuracy_score(test_labels, predictions)
```

Out[9]: 0.8349999999999996

```
In [10]: def get_meshgrid(data, step=.05, border=.5,):
    x_min, x_max = data[:, 0].min() - border, data[:, 0].max() + border
    y_min, y_max = data[:, 1].min() - border, data[:, 1].max() + border
    return np.meshgrid(np.arange(x_min, x_max, step), np.arange(y_min, y_max, step))

In [11]: def plot_decision_surface(estimator, train_data, train_labels, test_data, test_labels,
    colors = colors, light_colors = light_colors):

    #fit model
    estimator.fit(train_data, train_labels)


    #set figure size
    pyplot.figure(figsize = (16, 6))

    #plot decision surface on the train data
    pyplot.subplot(1,2,1)
    xx, yy = get_meshgrid(train_data)
    mesh_predictions = np.array(estimator.predict(np.c_[xx.ravel(), yy.ravel()])).reshape(xx.shape)
    pyplot.pcolormesh(xx, yy, mesh_predictions, cmap = light_colors)
    pyplot.scatter(train_data[:, 0], train_data[:, 1], c = train_labels, s = 100, cmap = colors)
    pyplot.title('Train data, accuracy={:.2f}'.format(metrics.accuracy_score(train_labels, mesh_predictions)))


    #plot decision surface on the test data
    pyplot.subplot(1,2,2)
    pyplot.pcolormesh(xx, yy, mesh_predictions, cmap = light_colors)
    pyplot.scatter(test_data[:, 0], test_data[:, 1], c = test_labels, s = 100, cmap = colors)
    pyplot.title('Test data, accuracy={:.2f}'.format(metrics.accuracy_score(test_labels, mesh_predictions)))

In [12]: def plot_decision_surface_neighbors(k):
    estimator = neighbors.KNeighborsClassifier(n_neighbors=k)
    plot_decision_surface(estimator, train_data, train_labels, test_data, test_labels)


In [13]: plot_decision_surface_neighbors(1)
```





```
In [14]: plot_decision_surface_neighbors(5)
```





```
In [15]: plot_decision_surface_neighbors(15)
```





```
In [16]: plot_decision_surface_neighbors(32)
```



In [17]: `plot_decision_surface_neighbors(120)`



In [18]: `plot_decision_surface_neighbors(500)`



```


In [19]: MAX_NEIGHBORS = 100
accuracy = []
for k in range(1, MAX_NEIGHBORS + 1):
    scores = []
    for train_indices, test_indices in model_selection.KFold(n_splits=5).split(classification_problem[0]):
        train_data = classification_problem[0][train_indices]
        train_labels = classification_problem[1][train_indices]
        test_data = classification_problem[0][test_indices]
        test_labels = classification_problem[1][test_indices]

        estimator = neighbors.KNeighborsClassifier(n_neighbors=k)
        estimator.fit(train_data, train_labels)
        predictions = estimator.predict(test_data)
        scores.append(metrics.accuracy_score(test_labels, predictions))
    accuracy.append(np.mean(scores))

In [20]: plt.figure(figsize=(20, 10))
plt.plot(np.arange(1, MAX_NEIGHBORS + 1), accuracy)
plt.show()

print("the best K is", 1 + argmax(accuracy))

```



the best K is 32

### 1.1.2 2. Наивный байесовский классификатор

```
In [23]: from sklearn.model_selection import cross_val_score
         from sklearn.naive_bayes import BernoulliNB, MultinomialNB, GaussianNB
```

```
In [24]: data_digits = datasets.load_digits()
         data_breast_cancer = datasets.load_breast_cancer()
```

```
In [25]: print(data_digits['data'])
         print(data_breast_cancer['data'])
```

```
[[ 0.  0.  5. ...,  0.  0.  0.]
 [ 0.  0.  0. ..., 10.  0.  0.]
 [ 0.  0.  0. ..., 16.  9.  0.]
 ...,
 [ 0.  0.  1. ...,  6.  0.  0.]
 [ 0.  0.  2. ..., 12.  0.  0.]
 [ 0.  0. 10. ..., 12.  1.  0.]]
[[ 1.79900000e+01  1.03800000e+01  1.22800000e+02 ...,  2.65400000e-01
  4.60100000e-01  1.18900000e-01]
 [ 2.05700000e+01  1.77700000e+01  1.32900000e+02 ...,  1.86000000e-01
  2.75000000e-01  8.90200000e-02]
 [ 1.96900000e+01  2.12500000e+01  1.30000000e+02 ...,  2.43000000e-01
  3.61300000e-01  8.75800000e-02]
 ...,
 [ 1.66000000e+01  2.80800000e+01  1.08300000e+02 ...,  1.41800000e-01
  2.21800000e-01  7.82000000e-02]
 [ 2.06000000e+01  2.93300000e+01  1.40100000e+02 ...,  2.65000000e-01
  4.08700000e-01  1.24000000e-01]
 [ 7.76000000e+00  2.45400000e+01  4.79200000e+01 ...,  0.00000000e+00
  2.87100000e-01  7.03900000e-02]]
```

digits - целые неотрицательные признаки

breast\_cancer - вещественные признаки

```
In [26]: for data in (data_digits, data_breast_cancer):
         for estimator in (BernoulliNB(), MultinomialNB(), GaussianNB()):
             print(cross_val_score(estimator=estimator, X=data.data, y=data.target).mean())
         print()
```

```
0.825823650778
0.870877148974
0.818600380355
```

```
0.627420402859
0.894579040193
0.936749280609
```

1. Каким получилось максимальное качество классификации на датасете breast\_cancer?  
94%
2. Каким получилось максимальное качество классификации на датасете digits?  
87%
3. Какие утверждения из приведенных ниже верны?
  - (a) На вещественных признаках лучше всего сработал наивный байесовский классификатор с распределением Бернулли

Нет, хуже

- (b) На вещественных признаках лучше всего сработал наивный байесовский классификатор с мультиномиальным распределением

Нет, с Гауссовым распределением

- (c) Мультиномиальное распределение лучше показало себя на выборке с целыми неотрицательными значениями признаков

Да, лучше, чем Гауссово и Бернулли. Нет, хуже чем с вещественными признаками

- (d) На вещественных признаках лучше всего сработало нормальное распределение

Да

### 1.1.3 3. Метрики в задаче регрессии


```
In [62]: from scipy import stats as sps
         from scipy.optimize import minimize
```

```
In [74]: MIN = -10
         MAX = 10
```

```
x = sps.uniform.rvs(MIN, MAX - MIN, size=500)
y = 0.5 * x + 1 + sps.norm.rvs(scale=0.2 ** 0.5, size=500)
```

```
In [75]: plt.figure(figsize=(20, 10))
         plt.scatter(x, y)
         plt.show()
```





```
In [76]: def function_mse(args):
          k, b = args
          return ((k * x + b - y) ** 2).mean()

          def function_mae(args):
              k, b = args
              return np.abs(k * x + b - y).mean()

In [77]: mse = minimize(function_mse, [0, 0])
          mae = minimize(function_mae, [0, 0])
          print(mse)
          print()
          print(mae)

          fun: 0.22413098492403824
          hess_inv: array([[ 0.01555316,  0.00479266],
                           [ 0.00479266,  1.00393264]])
          jac: array([ 1.63912773e-07,  2.42143869e-08])
          message: 'Optimization terminated successfully.'
          nfev: 20
          nit: 2
          njev: 5
          status: 0
          success: True
          x: array([ 0.4999781 ,  1.02264583])

          fun: 0.37848766814651036
          hess_inv: array([[ 5.98041558e-06,  5.36401568e-05],
                           [ 5.36401568e-05,  3.24649297e-01]])
```

```

jac: array([ 0.00247844,  0.          ])
message: 'Desired error not necessarily achieved due to precision loss.'
nfev: 451
nit: 10
njev: 110
status: 2
success: False
x: array([ 0.50240869,  1.01370381])

```

```


In [78]: k_mse, b_mse = mse.x
         k_mae, b_mae = mae.x

```

```

plt.figure(figsize=(20, 10))
plt.scatter(x, y)
plt.plot(x, k_mse * x + b_mse, label="MSE")
plt.plot(x, k_mae * x + b_mae, label="MAE")
plt.legend()
plt.show()

```



```

In [80]: new_x = np.append(x, sps.uniform.rvs(MIN, MAX - MIN, size=75))
         new_y = np.append(y, -1 + sps.norm.rvs(scale=0.2 ** 0.5, size=75))

```

```

In [81]: def function_mse(args):
         k, b = args
         return ((k * new_x + b - new_y) ** 2).mean()

         def function_mae(args):
         k, b = args
         return np.abs(k * new_x + b - new_y).mean()

```

```

In [82]: mse = minimize(function_mse, [0, 0])
         mae = minimize(function_mae, [0, 0])
         print(mse)
         print()
         print(mae)


         fun: 1.5886404915855898
hess_inv: array([[ 0.01538811,  0.00626163],
                 [ 0.00626163,  1.00294681]])
         jac: array([ 1.93715096e-07, -1.49011612e-08])
message: 'Optimization terminated successfully.'
         nfev: 20
         nit: 2
         njev: 5
         status: 0
         success: True
         x: array([ 0.43142489,  0.77433421])

         fun: 0.7068464887415729
hess_inv: array([[ 0.00054595, -0.00166921],
                 [-0.00166921,  0.00764435]])
         jac: array([ 0.00183568,  0.00104306])
message: 'Desired error not necessarily achieved due to precision loss.'
         nfev: 384
         nit: 9
         njev: 93
         status: 2
         success: False
         x: array([ 0.49134287,  0.99685401])

In [83]: k_mse, b_mse = mse.x
         k_mae, b_mae = mae.x

         plt.figure(figsize=(20, 10))
         plt.scatter(new_x, new_y)
         plt.plot(new_x, k_mse * new_x + b_mse, label="MSE")
         plt.plot(new_x, k_mae * new_x + b_mae, label="MAE")
         plt.legend()
         plt.show()

```



Видим, что МАЕ более устойчива к выбросам.

In [ ]: